

RL Traffic Signal Control

David Sanwald

May 8, 2018

Table of contents

1. Introduction

2. Reinforcement Learning

Introduction

Reinforcement Learning

Machine Learning

- Supervised Learning
 - Classification
 - Regression
- Unsupervised Learning
 - Clustering
 - ...
- Reinforcement Learning

Agent Environment




Figure 1: Agent environment interface

Markov Decision Process

MARKOV DECISION PROCESS is defined by quatuple $\langle \mathcal{S}, \mathcal{P}, \mathcal{R}, \mathcal{A} \rangle$

- \mathcal{S} , a set of states
- \mathcal{P} , a state transition matrix defining the probabilities of some possible next state s' given any state s
 $\mathcal{P}_{ss'}^a = \mathbb{P}[S_{t+1} = s' | S_t = s, A_t = a]$
- a reward function $\mathcal{R}_s^a = \mathbb{E}[R_{t+1} | S_t = s, A_t = a]$
- \mathcal{A} , a set of actions

- specifies agent's behaviour
- mapping of state to action

$$\pi(s) = a$$


$$\mathbb{P}(a|s) = \pi(a|s)$$

Markov Property

- The future is conditionally independent of the past given the presence
- implies memorylessness

$$\mathbb{P}[S_{t+1}|S_1, \dots, S_t] = \mathbb{P}[S_{t+1}|S_t]$$

Taxonomy of RL



Value Function

Expected return

- from state s and action a
- given policy π

$$Q^\pi(s, a) = \mathbb{E}[r_{t+1} + \gamma r_{t+2} + \gamma^2 r_{t+3} + \dots | s, a]$$

- decomposable into

$$Q^\pi(s, a) = \mathbb{E}[r + \gamma Q^\pi(s', a') | s, a]$$

Optimal Value Function

- optimal value function

$$Q^*(s, a) = \max_{\pi} Q^{\pi}(s, a) = Q^{\pi^*}(s, a)$$

- optimal policy

$$\pi^*(s) = \operatorname{argmax}_a Q^*(s, a)$$

- decomposition into

$$Q^*(s, a) = \mathbb{E}_{s'}[r + \gamma \max_{a'} Q^*(s', a') | s, a]$$

Off Policy learning

Q-learning

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha [R_{t+1} + \gamma \max_a Q(s_{t+1}, a) - Q(S_t, A_t)]$$

On Policy learning

Sarsa

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha [R_{t+1} + \gamma Q(s_{t+1}, A_{t+1}) - Q(S_t, A_t)]$$

Demo

Function Approximation

Why Function Approximation?

- large state spaces
- slow learning
- need for generalization

Naive Function Approximation

$$Q(s, a, \theta) \approx Q(s, a)$$

$$\mathcal{L}(\theta) = \mathbb{E} \left[\left(r + \gamma \max_{a'} Q(s', a', \theta) - Q(s, a, \theta) \right)^2 \right]$$

$$\frac{\partial \mathcal{L}(\theta)}{\partial \theta} = \mathbb{E} \left[\left(r + \gamma \max_{a'} Q(s', a', \theta) - Q(s, a, \theta) \right) \frac{\partial Q(s, a, \theta)}{\partial \theta} \right]$$

Deadly Triad


Deadly Triad

- function approximation
- off policy learning
- bootstrapping

Deadly Triad

- function approximation
- off policy learning
- bootstrapping

atari arcade games



HUMAN-LEVEL CONTROL THROUGH DEEP REINFORCEMENT LEARNING¹

- (almost) raw pixel input
- one agent/set of network weights
- comparable to human performance on 29 of 49 games

¹NATURE FEBRUARY 2015

experience replay

- decorrelates
- sample efficiency

target network

- inhibits loops

error clipping

- limits gradient magnitude

- store experience $e_t = (s_t, a_t, r_t, s_{t+1})$ in $D_t = \{e_1, \dots, e_t\}$
- at timestep t update $(s, a, r, s') \sim U(D)$

- separate target network $\tilde{Q}(s, a, \theta^-)$ and online network $Q(s, a, \theta)$
- TD error becomes $r + \gamma \max_{a'} Q(s', a', \theta^-) - Q(s, a, \theta)$

DQN architecture

